

Informational-computational trade-off in randomized numerical linear algebra

Xiaobo Yang

April 2023

Contents

1	Practical Applications	7
1.1	Datasets	7
1.2	Statistical applications	7
1.2.1	Gaussian process	7
2	Theory for application	9
2.1	Matrix-vector multiplication queries	9
2.1.1	Low rank approximation	9
2.1.2	Solve linear system	9
3	Theory for areas	11
3.1	Matrix-vector multiplication queries	11
3.1.1	Trace estimation	11
3.1.2	Eigenvalue estimation	14
4	Theoretic Frameworks	15
4.1	Meta-complexity	15
4.1.1	Gap-hamming distance	15
4.1.2	Spiked wishart matrix testing	17
5	Math Foundation	19
5.1	Key elements	19
5.1.1	Yao's principle	19
5.2	Concentration inequalities	19
5.2.1	Overlap of a vector on a large set	19
5.2.2	Hutchinson's trace estimator	20
5.2.3	Talagrand's inequality	20
5.3	Statistical distance	21
5.4	Random matrices	21

Preface

Explore information-theoretic lower bound and computational threshold(an upper bound yielded by algorithm) in randomized numerical linear algebra. We hope the gap is as small as possible.

Chapter 1

Practical Applications

Some coding and practical projects.

1.1 Datasets

xiaobo: wiki,arxiv matrix [Meyer et al., 2021]

1.2 Statistical applications

1.2.1 Gaussian process

xiaobo: train gaussian process in [Jiang et al., 2021] [Meyer et al., 2021]

Chapter 2

Theory for application

Specific complexity theory of specific applications in some areas.

2.1 Matrix-vector multiplication queries

2.1.1 Low rank approximation

[Bakshi et al., 2022]

xiaobo: reduced to estimation of minimum eigenvalue [Braverman et al., 2021]

2.1.2 Solve linear system

[Braverman et al., 2021]

xiaobo: reduced to estimation of minimum eigenvalue [Braverman et al., 2021]

Chapter 3

Theory for areas

General complexity theory of specific areas of problems.

3.1 Matrix-vector multiplication queries

We interact with the algorithm only through inputting query vectors and outputting matrix-vector product oracles.

3.1.1 Trace estimation

Problem 3.1.1 (trace estimation [Meyer et al., 2021]). *Given a matrix A , we input query vectors r_1, \dots, r_m and output product Ar_1, \dots, Ar_m to estimate $\text{tr}(A)$.*

There are a two types of query vectors.

1. (Adaptive) Vectors r_1, \dots, r_m are chosen adaptively. For example, $r_2 = Ar_1$. In this case, A may be used for several times.
2. (Non-adaptive) Vectors r_1, \dots, r_m are chosen independently. In this case, A can be used for only once, for example, streaming data.

Theorem 3.1.1 (Lower bound of adaptive queries [Meyer et al., 2021]). *Let A be a PSD matrix. If r_1, \dots, r_m are adaptive and with integer entries in $\{-2^b, \dots, 2^b\}$, then we need at least $m = \Omega\left(\frac{1}{\epsilon(b+\log(1/\epsilon))}\right)$ queries to output an estimate t so that, with probability at least $2/3$, $(1 - \epsilon)\text{tr}(A) \leq t \leq (1 + \epsilon)\text{tr}(A)$.*

Proof sketch for Theorem 3.1.1. First, reduce the problem to Gap-hamming distance problem(GHD) 4.1.1. For any vector $s, t \in \mathbb{R}^n$, reshape s, t to matrix $S, T \in \mathbb{R}^{n \times n}$. Let $Z = S + T$, $A = Z^T Z$.

Then, if we estimate $\text{tr}(A)$ with error $\epsilon = 1/\sqrt{n}$, we have

$$\begin{aligned} \widehat{\text{tr}(A)} &\geq 2(n + \sqrt{n})(1 - 1/\sqrt{n}) = 2n - 2 \quad \text{if } \langle s, t \rangle \geq \sqrt{n}, \\ \widehat{\text{tr}(A)} &\leq 2(n - \sqrt{n})(1 + 1/\sqrt{n}) = 2n - 2 \quad \text{if } \langle s, t \rangle \leq -\sqrt{n}, \end{aligned}$$

which means we can solve GHD problem by simply compare $\widehat{\text{tr}(A)}$ with $2n - 2$.

Finally, use lower bound of GHD. The reduction above cost $O(m\sqrt{n}(\log n + b)) = \Omega(n)$ bits by Theorem 4.1.1, thus we have $m = \Omega\left(\frac{1}{\epsilon(b+\log(1/\epsilon))}\right)$. \square

Theorem 3.1.2 (Lower bound of non-adaptive queries [Meyer et al., 2021]). *Let A be a PSD matrix. If r_1, \dots, r_m are non-adaptive, then we need at least $m = \Omega\left(\frac{1}{\epsilon}\right)$ queries to output an estimate t so that, with probability at least $3/4$, $(1 - \epsilon)\text{tr}(A) \leq t \leq (1 + \epsilon)\text{tr}(A)$.*

Proof sketch for Theorem 3.1.2. First, solving the problem implicitly solves a testing problem with m matrix-vector product oracles.

Problem 3.1.2 ([Meyer et al., 2021]). Fix $d, n \in \mathbb{N}$ such that $d \geq n$ and $n = 1/\epsilon$ for $\epsilon \in (0, 1]$. Let $A = G^T D G \in \mathbb{R}^{d \times d}$, where $G \in \mathbb{R}^{n \times d}$ generated by i.i.d $\mathcal{N}(0, 1)$ and $D = I_n := D_1$ or $\begin{bmatrix} I_{n-1} & 0 \\ 0 & 0 \end{bmatrix} := D_2$. Consider any algorithm taking a query matrix $U \in \mathbb{R}^{d \times m}$ as input and outputting product oracle $AU \in \mathbb{R}^{d \times m}$ to identify D .

By concentration inequality 5.2.2 of Hutchinson's estimator, we have that with probability at least $11/12$,

$$\frac{1}{d} \text{tr}(G^T D_1 G) \geq (1 - \epsilon/4) \text{tr}(D_1) \quad \text{and} \quad \frac{1}{d} \text{tr}(G^T D_2 G) \leq (1 + \epsilon/4) \text{tr}(D_2).$$

Then, if we estimate $t \approx \text{tr}(A)$ with probability at least $3/4$ and error $\epsilon/4$, we have

$$\begin{aligned} t &\geq (1 - \epsilon/4) \text{tr}(A) \geq (1 - \epsilon/4)^2 d/\epsilon \geq (1 - \epsilon/2) d/\epsilon \quad \text{if } D = D_1, \\ t &\leq (1 + \epsilon/4) \text{tr}(A) \leq (1 + \epsilon/4)^2 (1 - \epsilon) d/\epsilon \leq (1 - \epsilon/2) d/\epsilon \quad \text{if } D = D_2, \end{aligned}$$

which means we can identify D by simply compare t with $(1 - \epsilon/2) d/\epsilon$. By union bound, the probability is at least $2/3$.

Then, this problem can be further reduced to problem 4.1.2. Let $M = \begin{bmatrix} N^T N \\ L^T N \end{bmatrix}$, where $N \perp\!\!\!\perp L$ and $L \in \mathbb{R}^{n \times (d-m)}$ with i.i.d $\mathcal{N}(0, 1)$. Then we have $M \stackrel{\text{dist}}{=} G^T D G E_m$, where $E_m = [e_1, \dots, e_m]$ is the first m standard basis vectors. Therefore, it suffices to run algorithm above and if the output $D = D_1$ we claim $C = I$ otherwise $C = I - zz^T$.

Finally, use lower bound theorem 4.1.2, we have $m = \Omega(\frac{1}{\epsilon})$. □

Remark 3.1.1. 1. Lower bound of adaptive case is smaller than non-adaptive case. It make senses since adaptive queries contain more information on A .

2. The difference of two bounds is mathematically because intermediate problem 3.1.2 requires oracles (columns of U) are non-adaptive. But GHD problem doesn't have any requirement on relationships between oracle r s.
3. This inspires us, to reduce a problem to get a lower bound, we must keep key structures of problems as many as possible.

For instance, although we can reduce the non-adaptive scenario to GHD, the resulting bound is significantly weaker. This is because we abandon the critical constraint of exclusively using independent oracles r s to tackle the problem, which renders it more challenging.

xiaobo:

1. check why are lower bounds only suitable for PSD matrix.

Later, [Jiang et al., 2021] improve the lower bound with error probability δ .

Theorem 3.1.3 (Lower bound of non-adaptive queries with δ -error [Jiang et al., 2021]). Let A be a PSD matrix. If r_1, \dots, r_m are non-adaptive, then we need at least $m = \Omega\left(\frac{\sqrt{\log(1/\delta)}}{\epsilon} + \frac{\log(1/\delta)}{\log \log(1/\delta)}\right)$ queries to output an estimate t so that, with probability at least $1 - \delta$, $(1 - \epsilon) \text{tr}(A) \leq t \leq (1 + \epsilon) \text{tr}(A)$.

Proof sketch for Theorem 3.1.3. The proof is divided into two parts according to whether ϵ is small or large.

First, for small $\epsilon = O(1/\sqrt{\log(1/\delta)})$ we have $m = \Omega\left(\frac{\sqrt{\log(1/\delta)}}{\epsilon}\right)$:

1. Use Yao's principle (theorem 5.1.1) we only need to find a hard distribution. Since the bound is required to satisfied by any size n of matrix $A \in \mathbb{R}^{n \times n}$, we consider large n . Use theorem 5.4.1 to quantify the remaining randomness after we take several matrix-vector product oracles, we can observe if the number

of queries is not so large, then the remaining randomness will yield large variance and thus large error probability.

2. Specifically, let $W = I_n + \frac{1}{2C\sqrt{n}}(G + G^T)$ where entries of G are *i.i.d.* $\mathcal{N}(0, 1)$. Then, W is PSD by theorem 5.4.2 and $\text{tr}(W) \leq 2n$ *w.h.p.* for large n , since $\text{tr}(G) \sim \mathcal{N}(0, n)$. Quantitatively, for any $\delta \in (0, 1)$, take $n = \Omega(\log(1/\delta))$ then $\text{tr}(W) \leq 2n$ with probability at least $1 - \delta/10$. Then, for any $\epsilon = O(\frac{\sqrt{\log(1/\delta)}}{n}) = O(1/\sqrt{\log(1/\delta)})$, if we just take $m = n/2 = O\left(\frac{\sqrt{\log(1/\delta)}}{\epsilon}\right)$, use theorem 5.4.1 we prove the best estimation precision any algorithm can achieve *w.h.p.* is at most $c\sqrt{\log(1/\delta)}$ for $c = \Omega(1)$, which is large than required precision ϵ and thus yield lower bound $m = \Omega\left(\frac{\sqrt{\log(1/\delta)}}{\epsilon}\right)$.

Then, for any $\epsilon \in (0, 1)$, we have uniformly lower bound $m = \Omega\left(\frac{\log(1/\delta)}{\log \log(1/\delta)}\right)$:

1. The problem is further reduced to a testing problem 4.1.3. The reduction is naturally to consider trace estimation of a random matrix B generated from P or Q and find a gap between both confidence intervals. Actually, by concentration theorem 5.4.2 we have $W + 6\sqrt{\log(1/\delta)}I_n$ is PSD *w.h.p.*, and $|\text{tr}(W)| \leq 2\sqrt{2}\log(1/\delta)$ since $\text{tr}(W) \sim \mathcal{N}(0, 4\log(1/\delta))$. Then, for any non-adaptive trace estimation algorithm with precision ϵ , consider constant $C > \frac{10(1+\epsilon)}{1-\epsilon} - 6$, then for trace estimation t , we have
 - if $B \sim P$, then $t \geq (1-\epsilon)\text{tr}(B) = (1-\epsilon)((C+6)\log^{3/2}(1/\delta) - 2\sqrt{2}\log(1/\delta)) > 6(1+\epsilon)\log^{3/2}(1/\delta)$ with probability at least $1 - 2\delta$,
 - if $B \sim Q$, then $t \leq (1+\epsilon)\text{tr}(B) \leq (1+\epsilon)(6\log^{3/2}(1/\delta) + 2\sqrt{2}\log(1/\delta)) < 6(1+\epsilon)\log^{3/2}(1/\delta)$ with probability at least $1 - 2\delta$,

which distinguishes P and Q .

2. Use hardness theorem 4.1.3 we require at least $m = \Omega\left(\frac{\log(1/\delta)}{\log \log(1/\delta)}\right)$ queries.

Finally, combining these two bound by simple summation, we have our final lower bound. \square

Remark 3.1.2. Add spiked component $C\log^{3/2}(1/\delta)$ in the testing problem may be a standardized trick to analyze information-theoretic lower bound. Tuning the magnitude of the spiked component, we can show the phase transition by information-theoretic toolkit.

In this problem, the gap between informational and computational threshold is small(nearly matching).

Algorithm 1 Hutch++: randomized trace estimation with **adaptive** matrix-vector queries

Input: Matrix-vector product oracle for PSD matrix $A \in \mathbb{R}^{n \times n}$. Number m of queries.

Output: Approximation to $\text{tr}(A)$.

- 1: Sample $S \in \mathbb{R}^{n \times m/3}, G \in \mathbb{R}^{n \times m/3}$ with *i.i.d.* $\mathcal{N}(0, 1)$ entries.
 - 2: Compute an orthonormal basis $Q \in \mathbb{R}^{n \times m/3}$ for the span of AS via QR decomposition.
 - 3: **return** $t = \text{tr}(Q^T A Q) + \frac{3}{m}\text{tr}(G^T(I - QQ^T)A(I - QQ^T)G)$.
-

Theorem 3.1.4 (Upper bound of adaptive queries [Meyer et al., 2021]). Let $m = O\left(\frac{\sqrt{\log(1/\delta)}}{\epsilon} + \log(1/\delta)\right)$. Then with probability at least $1 - \delta$, Hutch++ yields an estimation t such that $(1-\epsilon)\text{tr}(A) \leq t \leq (1+\epsilon)\text{tr}(A)$.

Proof sketch for Theorem 3.1.4. First, \square

Algorithm 2 NA-Hutch++: randomized trace estimation with **non-adaptive** matrix-vector queries

Input: Matrix-vector product oracle for PSD matrix $A \in \mathbb{R}^{n \times n}$. Number m of queries.

Output: Approximation to $\text{tr}(A)$.

- 1: Fix constants c_1, c_2, c_3 such that $c_1 < c_2$ and $c_1 + c_2 + c_3 = 1$.
- 2: Sample $S \in \mathbb{R}^{n \times c_1 m}, R \in \mathbb{R}^{n \times c_2 m}, G \in \mathbb{R}^{n \times c_3 m}$ with *i.i.d.* $\mathcal{N}(0, 1)$ entries.
- 3: Let $Z = AR, W = AS$

4: **return** $t = \text{tr}((S^T Z)^+ (W^T Z)) + \frac{1}{c_3 m} (\text{tr}(G^T A G) - \text{tr}(G^T Z (S^T Z)^+ W^T G))$.

Theorem 3.1.5 (Upper bound of non-adaptive queries [Jiang et al., 2021]). *Let $m = O\left(\frac{\sqrt{\log(1/\delta)}}{\epsilon} + \log(1/\delta)\right)$. Then with probability at least $1 - \delta$, NA-Hutch++ yields an estimation t such that $(1 - \epsilon)\text{tr}(A) \leq t \leq (1 + \epsilon)\text{tr}(A)$.*

Proof sketch for Theorem 3.1.5.

□

3.1.2 Eigenvalue estimation

[Braverman et al., 2021]

Chapter 4

Theoretic Frameworks

General complexity frameworks for general theoretic areas.

4.1 Meta-complexity

Some classical problems that are usually reduced to.

4.1.1 Gap-hamming distance

Problem 4.1.1 (Gap-Hamming [Meyer et al., 2021]). *Let Alice and Bob be communicating parties who hold vectors $s \in \{\pm 1\}^n$ and $t \in \{\pm 1\}^n$, respectively. The Gap-Hamming problem asks Alice and Bob to return*

$$1 \text{ if } \langle s, t \rangle \geq \sqrt{n} \quad \text{and} \quad -1 \text{ if } \langle s, t \rangle \leq -\sqrt{n}.$$

In general, we need to compute the function $f(x, y)$ for $x \in \mathcal{X}, y \in \mathcal{Y}$, where \mathcal{X} and \mathcal{Y} are separate sources and we cannot simultaneously obtain both x and y . Instead, we repeatedly communicate information from x to y or from y to x , accumulating intermediate information during this process. Finally, we use this accumulative intermediate information to compute $f(x, y)$. The goal is to minimize the number of times of communication or the total amount of communication bits.

To formalize this communication process, a binary tree structure called a "protocol tree" T (Fig. 4.1) is used. Each possible value of $f(x, y)$ corresponds to a leaf node L in T . The communication complexity is measured by the height $H(T)$ of the tree T . Notably, for any node N in T , the set $R(N) := (x, y) : (x, y) \text{ can reach node } N \text{ always}$ forms a square, i.e., $R(N) = A \times B$ for some $A \subseteq \mathcal{X}$ and $B \subseteq \mathcal{Y}$. As a result, the leaf nodes form a rectangular partition of $\mathcal{X} \times \mathcal{Y}$.

By using the minimum number of square partitions M_f of the function $f(x, y)$, we can derive a lower bound on the communication complexity:

$$H(T) \geq \log_2 M_f.$$

Theorem 4.1.1 ([Chakrabarti and Regev, 2012] [Vidick, 2012] [Sherstov, 2012]). *Any randomized protocol solving problem 4.1.1 with probability $2/3$ needs at least $\Omega(n)$ bits of communication.*

Proof sketch for Theorem 4.1.1. The original proof is in [Chakrabarti and Regev, 2012] and [Vidick, 2012] make a simpler one. Here we use proof in [Sherstov, 2012], which is simplest among them.

First, use **Yao's principle** (theorem 5.1.1) to reduce the randomness of algorithm to randomness of a distribution μ of input (x, y) . Formally, consider all those deterministic protocol tree $T(x, y)$ such that $\mathbb{P}_{(x, y) \sim \mu}[T(x, y) \neq f(x, y)] \leq \epsilon$. Denote the minimum complexity (tree height) of such deterministic T by




Figure 1.1: A protocol tree

	y	y'	y''	y'''
x	0	0	0	1
x'	0	0	0	1
x''	0	0	0	0
x'''	0	1	1	1

Figure 1.2: The function f computed by the protocol of Figure 1.1

Figure 4.1: Protocol tree [Kushilevitz and Nisan, 1997]

$R_{\mu,\epsilon}(f)$, and denote minimum expectation complexity of random tree with error ϵ by $C_\epsilon(f)$, we have $C_\epsilon(f) \gtrsim R_{\mu,\epsilon}(f)$. See [Yao, 1983] lemma 2 in sect. 3.3 for details.

Then, it suffices to find a hard distribution μ of (x, y) and lower bound of $R_{\mu,\epsilon}(f)$. This is done by a trick called **ϵ -corruption bound**:

Lemma 4.1.1 (Yao's corruption bound [Sherstov, 2012]). *If $\exists \epsilon, \delta \in (0, 1)$ such that μ satisfies*

$$\mu(R \cap f^{-1}(+1)) > \epsilon \mu(R \cap f^{-1}(-1))$$

for every rectangle $R \subseteq X \times Y$ with $\mu(R) > \delta$, then $R_{\mu,\xi}(f) \geq \log_2 \left(\frac{1}{\delta} \left(\mu(f^{-1}(-1)) - \frac{\xi}{\epsilon} \right) \right)$ for small $\xi > 0$.

Intuitively, corruption guarantee the size of error area in each rectangle is not so small. Therefore, when total error is a given constant, we can lower bound the number of rectangle partition and thus lower bound the tree height. See [Yao, 1983] lemma 3 in sect. 3.3 for originality and details. Usually, people try to directly prove corruption bound for uniformly distribution.

Finally, in our setting, the Gap-Hamming distance problem(GHD_n) is further reduce to Gap-orthogonality problem(ORT_n) with $R_{1/3}(GHD_n) = \Omega(R_{1/3}(ORT_n))$, where $f(x, y) = (-1)\mathbf{1}_{|\langle x, y \rangle| \leq \sqrt{n}} + (+1)\mathbf{1}_{|\langle x, y \rangle| \geq 2\sqrt{n}}$. Note that $f^{-1}(+1) = \{(x, y) \mid |\langle x, y \rangle| \geq 2\sqrt{n}\}$, thus the corruption bound intuitively(not so rigorously) means for rectangle R with large probability, we have

$$\exists c, C > 0, \quad \mathbb{P}_{(x,y) \sim \text{Unif}[R]}[|\langle x, y \rangle| > C] \geq c, \quad (4.1)$$

which is essentially a **anti-concentration analysis**. Choosing uniform distribution is enough, it is proved by **probabilistic method** with Talagrand's inequalities(corollary 5.2.1) and a interesting combinatorial accounting trick:

Lemma 4.1.2 (Corruption bound [Sherstov, 2012]). *Let μ be uniform distribution on $\{\pm 1\}^n \times \{\pm 1\}^n$ and $R = A \times B$ be a rectangle such that $\mu(R \cap f^{-1}(+1)) \leq \epsilon \mu(R)$ then $\mu(R) = e^{-\Omega(n)}$. is small. In other word, if $\mu(R)$ is large, we have corruption bound.*

Proof sketch for lemma 4.1.2. The anti-concentration inequality is reduced to that, by corollary 5.2.1 we can find a set I of near-orthogonal vectors of $x \in A$ in rectangle $R = A \times B$ with large probability, and then by

some linear algebra and Hoffman-Wielandt inequality, for most of $y \in B$, there must exists a $x' \in I$ such that $|\langle x', y \rangle|$ is large. \square

Use the lemma, we compute $C_{1/3}(GHD_n) = \Omega(R_{1/3}(ORT_n)) = \Omega(n)$. \square

Remark 4.1.1. For general communication problem, the analysis of average-case complexity usually follows the roadmap above:

1. Use Yao's principle to reduce the problem with randomized algorithm into a randomized input with deterministic algorithm.
2. Find the math structures of complexity in setting of deterministic algorithm (In communication problem, it's height of a binary tree whose leaf nodes are rectangle partition).
3. Abstract these structures into pure math language, and then use fundamental math tool to work it out (In this communication function $f = GHD_n$, the structure is an anti-concentration analysis and tools are some probabilistic methods).

Remark 4.1.2. Another proof [Vidick, 2012] use a different description of the anti-concentration (theorem 5.2.1) in corruption bound (4.1), where they consider general vectors $x, y \in \mathbb{R}^n$ ($\mathcal{X}, \mathcal{Y}, R$ may be continuous).

4.1.2 Spiked wishart matrix testing

Problem 4.1.2 (Spiked wishart matrix [Meyer et al., 2021] [Perry et al., 2018]). Let $n = 1/\epsilon$ and let $z \in \mathbb{R}^n$ be a uniformly random unit vector. Let $N \in \mathbb{R}^{n \times m}$ be m i.i.d Gaussian vector drawn from an n -dimensional $\mathcal{N}(0, C)$, where $C = I$ or $I - zz^T$. Use N to identify C .

Theorem 4.1.2 ([Meyer et al., 2021]). To solve problem 4.1.2 with probability at least $2/3$, we need at least $m = \Omega(\frac{1}{\epsilon})$.

Proof sketch for Theorem 4.1.2. First, let P denote the distribution of N under null hypothesis $C = I$, and Q denote the distribution of N under alternative hypothesis $C = I - zz^T$. Any testing statistics ϕ outputs 1 for Q and 0 for P . It suffices to bound the total variation distance $d_{TV}(P, Q)$, since its control the summation of Type-1 and Type-2 errors

$$\min_{\phi} \{P(\phi = 1) + Q(\phi = 0)\} = 1 - d_{TV}(P, Q).$$

Then, by Pinsker's inequality, we have $d_{TV}(Q, P) \leq \sqrt{\frac{1}{2}D_{KL}(Q||P)} \leq \sqrt{\frac{1}{2}D_{\chi^2}(Q||P)}$. Therefore it suffices to upper bound $D_{\chi^2}(Q||P) = \int_N \left(\frac{Q(N)}{P(N)}\right)^2 P(N) dN - 1$. By theorem 5.3.1 with $\beta = -1$ and spiked prior z , we have

$$D_{\chi^2}(Q||P) = \mathbb{E}_{v, v'} \left[(1 - \langle v, v' \rangle^2)^{-m/2} \right] - 1,$$

where $v, v' \stackrel{\text{dist}}{=} z$ are uniformly random unit vectors. Therefore, by classical results of random unit vectors, p.d.f. of $x := \langle v, v' \rangle$ is $p(x) = \frac{\Gamma(n-1)}{2\Gamma((n-1)/2)^2} \left(\frac{1-x^2}{4}\right)^{(n-1)/2-1}$.

Finally, direct calculation yields, if $m = O(1/\epsilon)$, $\mathbb{E}_{v, v'} \left[(1 - \langle v, v' \rangle^2)^{-m/2} \right] < 6/5$ and thus $d_{TV}(Q, P) < 1/3$. Therefore, one of Type-1 and Type-2 errors must be larger than $1/3$. \square

Problem 4.1.3 ([Jiang et al., 2021]). Given $\delta \in (0, 1/2)$, set $n = \log(1/\delta)$. Independently take $g \sim \mathcal{N}(0, I_n)$ and $G \in \mathbb{R}^{n \times n}$ with i.i.d. $\mathcal{N}(0, 1)$ entries. Let $W = G + G^T$. Consider two distributions:

- Distribution P : $C \log^{3/2}(1/\delta) \cdot \frac{1}{\|g\|_2^2} gg^T + W + 6\sqrt{\log(1/\delta)} I_n$ for some fixed constant $C > 0$.
- Distribution Q : $W + 6\sqrt{\log(1/\delta)} I_n$.

Use **non-adaptive** queries q_1, \dots, q_m and oracles Aq_1, \dots, Aq_m to distinguish P and Q . In other word, take a matrix $Q \in \mathbb{R}^{n \times m}$ as query and AQ as oracle.

Theorem 4.1.3 ([Jiang et al., 2021]). *Any randomized algorithm solving problem 4.1.3 with probability $1 - \delta$ need at least $m = \Omega\left(\frac{\log(1/\delta)}{\log \log(1/\delta)}\right)$ queries.*

Proof sketch for Theorem 4.1.3. Similar to the proof of theorem 4.1.2, it is naturally to consider bounding the total variation $d_{TV}(P, Q)$, and by rotational invariance we WLOG assume $Q = E_m$, the first m standard basis vectors.

First, denote by P', Q' the distribution of $BQ(B \sim P$ or Q respectively). Let $L_{P'}, L_{Q'} \in \mathbb{R}^l$ be vectorization of matrices from P', Q' (remove the redundant variable, since B is symmetric), where $l = n + (n-1) + \dots + (n-m+1)$. Observe that conditioned on a realization g we have

$$d_{KL}(P', Q'|g) \leq d_{KL}(L_{P'}, L_{Q'}|g) \leq \|\mathbb{E}L_{P'} - \mathbb{E}L_{Q'}\|_2^2 = \sum_{i=1}^m \left\| C \log^{3/2}(1/\delta) \frac{gg^T}{\|g\|_2^2} e_i \right\|_2^2 = C^2 \log^3(1/\delta) \frac{\|g^T Q\|_2^2}{\|g\|_2^2}.$$

The first inequality use data processing inequality(theorem 5.3.3) and the second inequality uses theorem 5.3.2.

Then, find a typical event \mathcal{E} with positive probability. We take $\mathcal{E} = \left\{ \frac{\|g^T Q\|_2^2}{\|g\|_2^2} \leq \frac{1}{50C^2 n^3} \right\}$ and show that $\mathbb{P}(\mathcal{E}) \geq 10\delta$ if we only take $m = O\left(\frac{\log(1/\delta)}{\log \log(1/\delta)}\right)$ queries. Indeed, assume $m \leq n/2$

$$\mathbb{P}(\mathcal{E}) \geq \mathbb{P}(\|g\|_2^2 \geq n/2 \|g^T Q\|_2^2 \geq 1/(100C^2 n^2)) \cdot \mathbb{P}(\|g^T Q\|_2^2 \geq 1/(100C^2 n^2)) = \Omega(1) \cdot \Omega\left(\left(\frac{1}{n\sqrt{m}}\right)^m\right) = \Omega(e^{-\frac{m}{2} \log(n^2 m)}).$$

Therefore, when $m = O\left(\frac{\log(1/\delta)}{\log \log(1/\delta)}\right)$, $\mathbb{P}(\mathcal{E}) \geq 10\delta$.

Finally, conditioned realization $g \in \mathcal{E}$, we have the total variation $d_{TV}(P', Q'|g) \leq \sqrt{d_{TV}(P', Q'|g)/2} \leq 1/3$ by Pinsker's inequality. This means $\mathbb{P}[\text{the algorithms make mistake}|g] \geq 1/3$ under P' or Q' . Since $\mathbb{P}(\mathcal{E}) \geq 10\delta$, we have $\mathbb{P}[\text{the algorithms make mistake}] \geq \delta$ under P or Q if we only take $m = O\left(\frac{\log(1/\delta)}{\log \log(1/\delta)}\right)$ queries. \square

xiaobo: consider directly bound $d_{TV}(P, Q) \leq 1 - 2\delta$ instead?

Chapter 5

Math Foundation

Basic math tools.

5.1 Key elements

5.1.1 Yao's principle

When we analyze complexity of average-case problem with randomized algorithm, we can analyze a problem with randomized input and deterministic algorithm instead.

Theorem 5.1.1 (Yao's minimax principle [Yao, 1977]). *Given a problem \mathcal{P} , let \mathcal{A} be the set of all deterministic algorithms to solve \mathcal{P} and \mathcal{X} be the set of all instances of \mathcal{P} . Let $c(a, x)$ be the cost of algorithm $a \in \mathcal{A}$ solving instance $x \in \mathcal{X}$. For any distribution p over \mathcal{A} and any distribution q over input space \mathcal{X} , consider randomized algorithm $A \sim p$ and random instance $X \sim q$. Then, we have*

$$\max_{x \in \mathcal{X}} \mathbf{E}[c(A, x)] \geq \min_{a \in \mathcal{A}} \mathbf{E}[c(a, X)]$$

That is, the worst-case expected cost of the randomized algorithm is at least the expected cost of the best deterministic algorithm against input distribution q .

Proof for Theorem 5.1.1. Let $C = \max_{x \in \mathcal{X}} \mathbf{E}[c(A, x)]$ and $D = \min_{a \in \mathcal{A}} \mathbf{E}[c(a, X)]$. We have

$$C = \sum_x q_x C \geq \sum_x q_x \mathbf{E}[c(A, x)] = \sum_x q_x \sum_a p_a c(a, x) = \sum_a p_a \sum_x q_x c(a, x) = \sum_a p_a \mathbf{E}[c(a, X)] \geq \sum_a p_a D = D$$

□

5.2 Concentration inequalities

5.2.1 Overlap of a vector on a large set

If two subsets of \mathbb{R}^n are "large", then the "overlap" of them is large. Here, "large" means the subsets have large measure under Gaussian measure, and "overlap" means: For any non-empty $A, B \subseteq \mathbb{R}^n$. Denote by $\gamma_{|A \times B}$ the probability measure corresponding to the normalized restriction of $\mathcal{N}(0, I_n) \times \mathcal{N}(0, I_n)$ to $A \times B$ and let

$$v(A, B) := \mathbb{E}_{(x, y) \sim \gamma_{|A \times B}} [|\langle x, y \rangle|^2].$$

Formally, we have

Theorem 5.2.1 ([Vidick, 2012]). *For any $\eta > 0$, there exists¹ a $\delta > 0$ such that for all large enough n , if A, B both have measure $\gamma(A), \gamma(B) \geq e^{-\delta n}$ then*

$$v(A, B) \geq (1 - \eta) v(\mathbb{R}^n, \mathbb{R}^n) = (1 - \eta)n$$

It's continuous version of (4.1), which is more general to yield corruption bound for communication problem with continuous domain \mathcal{X}, \mathcal{Y} and rectangle R .

5.2.2 Hutchinson's trace estimator

A simple estimator of trace of matrix by monte carlo method and we use only matrix-vector product oracle.

Definition 5.2.1 (Hutchinson's Estimator). *Given a matrix $A \in \mathbb{R}^{d \times d}$. We estimate $\text{tr}(A)$ by*

$$H_m(A) = \frac{1}{m} \sum_{i=1}^m g_i^T A g_i = \frac{1}{m} \text{tr}(G^T A G),$$

where $G = [g_1, \dots, g_m] \in \mathbb{R}^{d \times m}$ is a Gaussian matrix with i.i.d. $\mathcal{N}(0, 1)$ entries.

Theorem 5.2.2 (Hutchinson analysis [Meyer et al., 2021]). *Let $A \in \mathbb{R}^{d \times d}$, $\delta \in (0, 1/2]$, $\ell \in \mathbb{N}$. Let $H_\ell(A)$ be the ℓ -query Hutchinson estimator defined in (1), implemented with mean 0, i.i.d. sub-Gaussian random variables with constant sub-Gaussian parameter. For fixed constants c, C , if $\ell > c \log(1/\delta)$, then with probability $\geq 1 - \delta$,*

$$|H_\ell(A) - \text{tr}(A)| \leq C \sqrt{\frac{\log(1/\delta)}{\ell}} \|A\|_F.$$

So, if $\ell = O\left(\frac{\log(1/\delta)}{\varepsilon^2}\right)$ then, with probability $\geq 1 - \delta$, $|H_\ell(A) - \text{tr}(A)| \leq \varepsilon \|A\|_F$.

Proof sketch for Theorem 5.2.2. First, vectorize gaussian matrix G to $\bar{g} \in \mathbb{R}^{dl}$ and let $\bar{A} = \text{diag}\{A, A, \dots, A\} \in \mathbb{R}^{dl \times dl}$.

Then, $H_\ell(A) = \bar{g}^T \bar{A} \bar{g} / \ell$ and use Hanson-Wright inequality [Vershynin, 2019]. \square

5.2.3 Talagrand's inequality

Theorem 5.2.3 (Talagrand [Sherstov, 2012]). *For a fixed convex set $S \subseteq \mathbb{R}^n$ and a random $x \in \{-1, +1\}^n$,*

$$\mathbb{P}[x \in S] \mathbb{P}[\rho(x, S) > t] \leq e^{-t^2/16},$$

where $\rho(x, S) = \inf_{y \in S} \|x - y\|$.

Intuitively, it means measure concentration around convex hull of large subset $S \subseteq \{\pm 1\}^n$, where the measure means counting measure (or uniform distribution). It can be seen as an isoperimetric inequality in discrete structures.

The usefulness of Talagrand's inequality is led by appropriately choosing set S , according to property we interest. Here are some useful corollaries.

Corollary 5.2.1 ([Sherstov, 2012]). *For every linear subspace $V \subseteq \mathbb{R}^n$ and every $t > 0$, one has*

$$\mathbb{P}_{x \in \{-1, +1\}^n} \left[\left| \|\text{proj}_V x\| - \sqrt{\dim V} \right| > t \right] < 4e^{-ct^2},$$

where $c > 0$ is an absolute constant.

Proof for Theorem 5.2.1. First, note that it suffices to prove

$$\mathbf{P}[\left| \rho(x, V) - \sqrt{n - \dim V} \right| > t] \leq 4e^{-\Omega(t^2)}$$

and $n - \dim V = \text{tr}(\mathbb{E}[x^T (I - P_V) x]) = \mathbb{E}[\rho(x, V)^2]$, where P_v is orthogonal projection on V . Therefore, take $S = \{v \in \mathbb{R}^n : \rho(v, V) \leq a\}$, then by Talagrand's inequality we have

$$\mathbb{P}[\rho(x, V) \leq a] \mathbb{P}[\rho(x, V) > a + t] \leq \mathbb{P}[\rho(x, V) \leq a] \mathbb{P}[\rho(x, S) > t] \leq e^{-t^2/16}.$$

Then, take $a = \text{median}(\rho(x, V)) := m$, we have

$$\begin{aligned} \mathbb{P}[\rho(x, V) > m + t] &\leq 2e^{-t^2/16}, \\ \mathbb{P}[\rho(x, V) \leq m - t] &\leq 2e^{-t^2/16}. \end{aligned}$$

Finally, the result follows $m = \sqrt{\mathbb{E}[\rho(x, V)^2]} + O(1)$. See [Tao, 2009] for details. \square

5.3 Statistical distance

Theorem 5.3.1 ([Perry et al., 2018] proposition 5.11). *For any $|\beta| < 1$, there exists $\delta > 0$ such that the following holds. Let $\mathcal{X} = \{\mathcal{X}_n\}$ be a family of prior distribution of spiked vectors x with $1 - \delta \leq \|x\| \leq 1 + \delta$. Let Q_n be joint distribution of N i.i.d samples from $\mathcal{N}(0, I_n + \beta x x^T)$ and P_n be joint distribution of N i.i.d samples from $\mathcal{N}(0, I_n)$. Then we have*

$$\mathbb{E}_{P_n} \left[\left(\frac{dQ_n}{dP_n} \right)^2 \right] = \mathbb{E}_{x, x' \sim \mathcal{X}_n} \left[(1 - \beta^2 \langle x, x' \rangle^2)^{-N/2} \right].$$

Proof for Theorem 5.3.1. First, expand LHS directly, we have $\mathbb{E}_{P_n} \left[\left(\frac{dQ_n}{dP_n} \right)^2 \right] = \mathbb{E}_{Q_n} \left[\frac{dQ_n}{dP_n} \right]$

$$\begin{aligned} \frac{dQ_n}{dP_n} (y_1, \dots, y_N) &= \mathbb{E}_{x' \sim \mathcal{X}} \left[\prod_{i=1}^n \frac{\exp \left(-\frac{1}{2} y_i^\top (I + \beta x' x'^\top)^{-1} y_i \right)}{\sqrt{\det(I + \beta x' x'^\top)} \exp \left(-\frac{1}{2} y_i^\top y_i \right)} \right] \\ &= \mathbb{E}_{x'} \left[\det(I + \beta x' x'^\top)^{-N/2} \prod_{i=1}^N \exp \left(-\frac{1}{2} y_i^\top \left((I + \beta x' x'^\top)^{-1} - I \right) y_i \right) \right]. \end{aligned}$$

Then, simplify it by Sherman–Morrison formula, we have $(I + \beta x' x'^\top)^{-1} - I = \frac{-\beta}{1 + \beta \|x'\|^2} x' x'^\top$, and then

$$= \mathbb{E}_{x'} \left[\left(1 + \beta \|x'\|^2 \right)^{-N/2} \prod_{i=1}^N \exp \left(\frac{1}{2} \frac{\beta}{1 + \beta \|x'\|^2} \langle y_i, x' \rangle^2 \right) \right].$$

Finally, passing to the second moment, we compute

$$\begin{aligned} \mathbb{E}_{P_n} \left[\left(\frac{dQ_n}{dP_n} \right)^2 \right] &= \mathbb{E}_{x, x'} \left[\left(1 + \beta \|x'\|^2 \right)^{-N/2} \prod_{i=1}^N \mathbb{E}_{y_i \sim \mathcal{N}(0, I + \beta x x^\top)} \exp \left(\frac{1}{2} \frac{\beta}{1 + \beta \|x'\|^2} \langle y_i, x' \rangle^2 \right) \right] \\ &= \mathbb{E}_{x, x'} \left[\left(1 + \beta \|x'\|^2 \right)^{-N/2} \prod_{i=1}^N \left(1 - \frac{\beta}{1 + \beta \|x'\|^2} \left(\|x'\|^2 + \beta \langle x, x' \rangle^2 \right) \right)^{-1/2} \right] \\ &= \mathbb{E}_{x, x'} \left[\left(1 - \beta^2 \langle x, x' \rangle^2 \right)^{-N/2} \right] \end{aligned}$$

Note that, the condition about δ is because here the MGF step requires $\frac{\beta}{1 + \beta \|x'\|^2} \left(\|x'\|^2 + \beta \langle x, x' \rangle^2 \right) < 1$. \square

Theorem 5.3.2 ([Jiang et al., 2021]). *For distribution $P = \mathcal{N}_k(\mu_1, \Sigma_1)$ and $Q = \mathcal{N}_k(\mu_2, \Sigma_2)$, we have*

$$d_{KL}(P, Q) = \frac{1}{2} \left\{ (\mu_2 - \mu_1)^T \Sigma_2^{-1} (\mu_2 - \mu_1) + \text{tr}(\Sigma_2^{-1} \Sigma_1) - \log \frac{\det(\Sigma_1)}{\det(\Sigma_2) - k} \right\}.$$

Theorem 5.3.3 (Data processing inequality [Jiang et al., 2021]). *For random variable X, Y and any function f , we have*

$$d_{KL}(f(X), f(Y)) \leq d_{KL}(X, Y).$$

5.4 Random matrices

In this section, we denote the distribution of random matrices $G \in \mathbb{R}^{n \times n}$ with i.i.d. $\mathcal{N}(0, 1)$ entries by $\mathcal{N}(n)$.

Theorem 5.4.1 (Remaining randomness [Simchowitz et al., 2020] [Jiang et al., 2021]). *Let $G \sim \mathcal{N}(n)$. Let $W = (G + G^T)/2$. For any sequence of vector queries v_1, \dots, v_T , along with oracles $w_i = W v_i$. Then, conditioned on these observations, there exists a rotation matrix U , independent of w_i , such that*

$$U W U^T = \begin{bmatrix} Y_1 & Y_2^T \\ Y_2 & \tilde{W} \end{bmatrix},$$

where Y_1, Y_2 are deterministic and $\tilde{W} = (\tilde{G} + \tilde{G}^T)/2$, where $\tilde{G} \sim \mathcal{N}(n - T)$.

Proof sketch for Theorem 5.4.1.

□

Theorem 5.4.2 (concentration of the largest singular value [Jiang et al., 2021]). *Let $G \sim \mathcal{N}(n)$. Then, for any $t \geq 0$ we have*

$$\mathbb{P}[s_{\max}(G) \leq 2\sqrt{n} + t] \geq 1 - 2e^{-t^2/2}.$$

Proof sketch for Theorem 5.4.2.

□

Bibliography

- [Bakshi et al., 2022] Bakshi, A., Clarkson, K. L., and Woodruff, D. P. (2022). Low-rank approximation with $1/\epsilon^{1/3}$ matrix-vector products.
- [Braverman et al., 2021] Braverman, M., Hazan, E., Simchowicz, M., and Woodworth, B. (2021). The gradient complexity of linear regression.
- [Chakrabarti and Regev, 2012] Chakrabarti, A. and Regev, O. (2012). An optimal lower bound on the communication complexity of gap-hamming-distance.
- [Jiang et al., 2021] Jiang, S., Pham, H., Woodruff, D. P., and Zhang, Q. R. (2021). Optimal sketching for trace estimation.
- [Kushilevitz and Nisan, 1997] Kushilevitz, E. and Nisan, N. (1997). *Communication Complexity*. Cambridge University Press, Cambridge.
- [Meyer et al., 2021] Meyer, R. A., Musco, C., Musco, C., and Woodruff, D. P. (2021). Hutch++: Optimal stochastic trace estimation.
- [Perry et al., 2018] Perry, A., Wein, A., Bandeira, A., and Moitra, A. (2018). Optimality and sub-optimality of PCA I: Spiked random matrix models. *Annals of Statistics*, 46:2416–2451.
- [Sherstov, 2012] Sherstov, A. A. (2012). The communication complexity of gap hamming distance.
- [Simchowicz et al., 2020] Simchowicz, M., Alaoui, A. E., and Recht, B. (2020). Tight query complexity lower bounds for pca via finite sample deformed wigner law.
- [Tao, 2009] Tao, T. (2009). Talagrand’s concentration inequality. Weblog entry.
- [Vershynin, 2019] Vershynin, R. (2019). *High-Dimensional Probability: An Introduction with Applications in Data Science*.
- [Vidick, 2012] Vidick, T. (2012). A concentration inequality for the overlap of a vector on a large set with application to the communication complexity of the gap-hamming-distance problem.
- [Yao, 1977] Yao, A. C.-C. (1977). Probabilistic computations: Toward a unified measure of complexity.
- [Yao, 1983] Yao, A. C.-C. (1983). Lower bounds by probabilistic arguments.